# 3D Computer Vision 

Visual Motion Estimation

Ivan Marković<br>University of Zagreb Faculty of Electrical Engineering and Computing Department of Control and Computer Engineering<br>Laboratory for Autonomous Systems and Mobile Robotics (lamor.fer.hr)

## Outline

- Problem formulation
-3D-3D motion estimation
-3D-2D motion estimation
- Perspective from $n$ points (PnP)
- Stereo and mono odometry using PnP
- 2D-2D motion estimation
- Estimating the essential matrix
- Relative scale
- Keyframe selection


## Outline

- Problem formulation
-3D-3D motion estimation
- 3D-2D motion estimation
- Perspective from n points (PnP)
- Stereo and mono odometry using PnP
- 2D-2D motion estimation
- Estimating the essential matrix
- Relative scale
- Keyframe selection


## VO - front-end vs. back-end



The part of visual odometry in charge of detecting and tracking (or matching) features and estimating the relative motion is called the front-end.

The part of visual odometry in charge of refining motion estimation and ensuring local consistency over a window of past frames is called the back-end.

The motion estimation block usually involves a robust model estimation procedure such as RANSAC that is used to determine the set of inliers based on which the final motion parameters are computed.

To estimate the motion between two views, we need to be able to detect and match features of one or several subsequent frames; thus, the following assumptions are made:

1. Adequate illumination of the scene - if the images are over- or undersaturated it will be nearly impossible to detect and track features.
2. Dominance of static objects over the moving objects - to estimate ego-motion we must use features on static parts of the scene; otherwise, we do not know if the scene is moving and the camera is static or vice-versa.
3. Enough texture to allow apparent motion to be estimated - moving in featurless environments is almost the same as moving in darkness
4. Sufficient scene overlap between subsequent frames - otherwise, we do not have any features to match

## Problem formulation

A camera is moving through an environment and taking images at discrete time steps $k$. It can be handheld or rigidly attached to a platform (such as a mobile robot or a vehicle).
In the case of a monocular system, we denote the set of $k$ images as

$$
I_{0: k}=\left\{I_{0}, I_{1}, \ldots, I_{k}\right\} .
$$

In the case of a stereo system, at each time step we have a left and a right image that we denote by

$$
\begin{aligned}
& I_{l, 0: k}=\left\{I_{l, 0}, I_{l, 1}, \ldots, I_{l, k}\right\} \\
& I_{r, 0: k}=\left\{I_{r, 0}, I_{r, 1}, \ldots, I_{r, k}\right\} .
\end{aligned}
$$

In the stereo case, we can use the coordinate system of the left camera as the origin frame.

Two camera positions at subsequent time steps $k-1$ and $k$ are related by the rigid body transformation

$$
T_{k}^{k-1}=\left[\begin{array}{cc}
R_{k}^{k-1} & t_{k}^{k-1} \\
0 & 1
\end{array}\right]
$$

The set $T_{0: k}=\left\{T_{1}^{0}, T_{2}^{1}, \ldots, T_{k}^{k-1}\right\}$ contains all subsequent motions.
The set of camera poses

$$
C_{0: k}=\left\{C_{0}, C_{1}, \ldots, C_{k}\right\}
$$

contains the transformations of the camera coordinate frame with respect to the initial coordinate frame at $k=0$.



The current pose of the camera $C_{k}$, or of the mobile robot or a vehicle, is computed by concatenating all the transformations from $0: k$

$$
C_{k}=C_{k-1} T_{k}=C_{k-2} T_{k-1} T_{k}=\cdots=C_{0} T_{1} T_{2} \ldots T_{k-1} T_{k}
$$

Note that for brevity we omit the previous frame index, i.e., $T_{k} \leftarrow T_{k}^{k-1}$.

The main task of visual odometry is to compute the relative rigid body transformations $T_{k}$ from images $I_{k}$ and $I_{k-1}$ and then to concatenate them in order to recover the full trajectory $C_{0: k}$ of the camera.
Evidently, visual odometry recovers the trajectory incrementally, pose after pose.
As discussed earlier, the optional step includes optimization over a local window of frames that aims to enforce local consistency and hopefully result with a more accurate estimate of the total trajectory.


Local optimization over $m$ frames

Depending whether point feature correspondence $p_{k}$ and $p_{k-1}$ are specified in 2D or 3D we have three groups of motion estimation approaches:

1. 2D-2D: both $p_{k}$ and $p_{k-1}$ are specified in 2D image coordinates
2. 3D-2D: point features $p_{k-1}$ are specified in 3D coordinates $X_{k-1}$ and $p_{k}$ are their corresponding 2D image reprojections on the image $I_{k}$.
3. 3D-3D: both $p_{k}$ and $p_{k-1}$ are specified in 3D coordinates $X_{k}$ and $X_{k-1}$, respectively. For this approach, we have to triangulate 3D points and each time step, e.g., by using a stereo or depth camera.
In the following, we discuss each of these approaches and their minimal case solutions.

## Outline

- Problem formulation
- 3D-3D motion estimation
-3D-2D motion estimation
- Perspective from n points (PnP)
- Stereo and mono odometry using PnP
- 2D-2D motion estimation
- Estimating the essential matrix
- Relative scale
- Keyframe selection


## 3D-3D motion estimation

The 3D-3D approach estimates motion from 3D-to-3D point feature correspondences - $X_{k}$ and $X_{k-1}$ - it is also called the point cloud registration problem.
The minimal-case solution involves 3 non-collinear correspondences.
In general, the goal is to find the aligning transformation $T_{k}$ that minimizes the following distance between the two 3D features sets

$$
\underset{T_{k}}{\arg \min } \sum_{i}\left\|X_{k}^{i}-T_{k} X_{k-1}^{i}\right\|^{2},
$$

where the superscript $i$ denotes the $i$-th point feature.

## 3D-3D motion estimation

A closed-form solution for registration of 3D-3D correspondences exists and is based on least square fitting?
Another class of solutions are the iterative closest points (ICP) algorithms that minimize point-to-point, point-to-line, or point-to-plane distances (sensitive to initial guess). An algorithm that joins multiple criteria in the minimization function is called the Generalized ICP algorithm².
Essentially, any point cloud registration approach discussed in earlier lectures can be used at this point, but keep in mind the final application - if its robotics or autonomous vehicles, then the algorithm should be executed in real-time, i.e., at least at the rate of 10 Hz .


[^0]${ }^{2}$ A. Segal, D. Hehnel, S. Thrun (2005). „Generalized-ICP."

## Aligning 3D point clouds = motion estimation

As stated earlier, to compute motion estimation from 3D-3D correspondences we triangulate the point features at steps $k-1$ and $k$. Thereafter, we align the 3D points using any available registration method.


## Aligning 3D point clouds = motion estimation

As stated earlier, to compute motion estimation from 3D-3D correspondences we triangulate the point features at steps $k-1$ and $k$. Thereafter, we align the 3D points using any available registration method.


## Aligning 3D point clouds = motion estimation

As stated earlier, to compute motion estimation from 3D-3D correspondences we triangulate the point features at steps $k-1$ and $k$. Thereafter, we align the 3D points using any available registration method.


## Aligning 3D point clouds = motion estimation

As stated earlier, to compute motion estimation from 3D-3D correspondences we triangulate the point features at steps $k-1$ and $k$. Thereafter, we align the 3D points using any available registration method.


## 3D-3D VO pipeline

The pipeline of a 3D-3D odometry might have the following structure:

1. Capture two stereo image pairs $I_{l, k-1}, I_{r, k-1}$ and $I_{l, k}, I_{r, k}$.
2. Detect and match features between $I_{l, k}$ and $I_{l, k-1}$.
3. Match and triangulate matched features for each stereo pair
4. Compute the relative transformation $T_{k}$ from triangulated 3D point feature sets $X_{k-1}$ and $X_{k}$.
5. Concatenate camera transformation by computing $C_{k}=C_{k-1} T_{k}$.

## Outline

- Problem formulation
-3D-3D motion estimation
- 3D-2D motion estimation
- Perspective from $n$ points (PnP)
- Stereo and mono odometry using PnP
- 2D-2D motion estimation
- Estimating the essential matrix
- Relative scale
- Keyframe selection


## 3D-2D motion estimation

The 3D-2D approach estimates motion from 3D-to-2D point feature correspondences - $X_{k-1}$ and $p_{k}$ - the reprojections of the 3D point features to the image $I_{k}$ - it is also called the perspective from $n$ points (PnP).
The minimal-case solution involves 3 non-collinear points ( +1 for disambiguation).
In general, the goal is to find the transformation $T_{k}$ that minimizes the image reprojection error

$$
\underset{T_{k}}{\arg \min } \sum_{i}\left\|p_{k}^{i}-\pi_{k}\left(T_{k} X_{k-1}^{i}\right)\right\|^{2},
$$

where the superscript $i$ denotes the $i$-th point feature and $\pi_{k}($.$) is the$ projection function that projects a 3D point to image frame $I_{k}$.

## Feature depth uncertainty

The 3D-2D approach has been found to yield more accurate results than 3D$3 D^{3}$ correspondences because it minimizes the image reprojection error instead of the 3D-3D feature position error.
The reason behind this lies in the fact that the uncertainty of feature 3D positions is highly anisotropic and the depth uncertainty increases the larger the ratio of object's distance with respect to cameras baseline.
The uncertainty of reprojections, on the other hand, is highly isotropic and it is less error sensitive to minimize the image reprojection error than alignment of triangulated 3D-3D
 point feature sets ${ }^{4}$.
${ }^{3}$ Nister, D., Naroditsky, O., and Bergen, J. (2004). „Visual odometry."
${ }^{4}$ Badino, H., Yamamoto, A., and Kanade, T. (2013). „Visual odometry by multi-frame feature integration."

## PnP problem

Perspective from $n$ points ( PnP ) is actually a camera localization problem, where we aim to estimate the 6DoF pose of the camera with respect to the world frame from a set of 3D-2D point feature correspondences.
Minimal-case solution:

1. 1 point - infinite solutions
2. 2 points - infinitely many solution but on a circular arc
3. 3 points (non-collinear) - up to 4 possible solutions
4. 4 points - unique solution.

## World

 frame

1 point and 2 points analysis

1 point - infinetly many solutions
2 points - bounded on a circular arc


## P3P method

P3P problem was first introduced and solved in 1841. It was discovered that this problem typically does not lead to a unique mathematical solution, but rather may yield up to four distinct solutions, any of which could be the pose of the camera ${ }^{5}$.
There are multiple methods for solving the P3P analytically and most are a two-step process:

1. Determine the length of projection rays
2. Estimate the camera pose.

We will explore the approach that systematically analyzes all the possible solutions and takes into account constraints that correspond to geometrically nondegenerate solutions ${ }^{6}$.


[^1]
## P3P method



Simply put, the goal of P3P is to find the coordinates of points $A, B, C$ in the camera coordinate frame

$$
A^{C}=R_{W}^{C} A^{W}+t_{W}^{C} .
$$

Note that positions of the points are known in the world frame and we assume to detect the point reprojections in the image.
Early approaches solved for positions of $A, B, C$ in the camera frame and then used a 3D point cloud registration method for the camera pose.

## P3P method



Direct approaches aim at solving the geometry of tetrahedron and start with the law of cosines to obtain set of polynomial eqs:

$$
\begin{aligned}
& s_{1}^{2}=L_{B}^{2}+L_{A}^{2}-2 L_{B} L_{A} \cos \theta_{a b} \\
& s_{2}^{2}=L_{A}^{2}+L_{C}^{2}-2 L_{A} L_{C} \cos \theta_{a c} \\
& s_{3}^{2}=L_{B}^{2}+L_{C}^{2}-2 L_{B} L_{C} \cos \theta_{b c}
\end{aligned}
$$

Note that $L_{A}, L_{B}, L_{C}$ represent unknowns - angles determined from the image points (the camera is calibrated) and triangle sides are known from the world frame 3D coordinates.

## P3P method

With some algebra we can reduce the problem to a set of two polynomial equation of the $2 n d$ order

$$
\begin{aligned}
& (1-u)^{2} y^{2}-u x^{2}-2 y \cos \theta_{b c}+2 x y u \cos \theta_{a b}+1 \\
& (1-v)^{2} x^{2}-v y^{2}-2 x \cos \theta_{a c}+2 x y v \cos \theta_{a b}+1
\end{aligned}
$$

where $x=L_{A} / L_{C}, y=L_{B} / L_{C}, u=s_{3}^{2} / s_{1}^{2}, v=s_{2}^{2} / s_{1}^{2}$.
For a system of $n$ polynomial equations in $n$ variables, the number of solutions is equal to the product of the equation degrees.
In our case, all unknowns are either linear or quadratic, thus we have 4 possible solutions, i.e., 4 possible combinations of $L_{A}, L_{B}, L_{C}$ that satisfy the equations - to disambiguate the solution a 4th point is usually used.

[^2]
## P3P ambiguity



For example, assume that triangle sides and and projection ray angles are equal.

By rotating the triangle along one of its sides the opposite vertice will at one point intersect the projection and an
equilateral triangle will be formed again.

An equivavlent effect can be also obtained my moving the camera's coordinate system.

The 4th point can be used to solve another P3P and see which of the 8 solutions overlap.

## Robust P3P

With the unique solution for $L_{A}, L_{B}, L_{C}$ coordinates in the camera frame, we can now compute the rigid body transformation between the camera and world frame, i.e., localize the camera, by using any of the existing 3D point cloud registration algorithms, e.g., the closed-form least squares fitting. In reality, we usually have much more than 3 points available and the correspondences that are not perfect. To obtain a robust solution in practice, we couple the P3P method with RANSAC:

1. Select 3 points randomly
2. Estimate the camera pose using P3P
3. Count the points that support this hypothesis (by comparing their reprojections from 3D position and detected points in the image)
4. Select the best solution as the final solution.

## One-stage solvers

The above-described method is a two-stage method: it first determines the lengths of projection rays and then obtains the rotation and translation via point alignment method.

The second stage usually involves matrix decomposition which is timeconsuming and highly sensitive to the distances obtained from the first step, thus reducing the efficiency and accuracy of the final solutions.
There exists approaches that directly solve for the rotation and translation parameters and are called one-stage solvers. They do not experience alignment issues, achieve higher speed and accuracy. They are usually based on clever parametrization and change of coordinate frames ${ }^{8}$.

## Using more than 3 points

All the aforementioned approaches are 3-point methods. Note that the RANSAC approach only enables us to find the set of inliers, but in the end, we still have only methods that return solutions based only on three points.
To estimate the camera pose from $n \geq 4$ points, e.g., the inlier set returned by RANSAC, we can use the EPnP algorithm ${ }^{9}$. The idea is to express the $n$ 3D points as a weighted sum of 4 virtual control points. The problem then reduces to estimating the coordinates of these control points in the camera frame, which can be done in $O(n)$.
The EPnP method expresses these coordinates as weighted sum of the eigenvectors of a $12 \times 12$ matrix and solves a small constant number of quadratic equations to pick the right weights.

[^3]
## Refinement via nonlinear optimization

Finally, note that the optimization problem that we were solving had the following form

$$
\underset{T_{k}}{\arg \min } \sum_{i}\left\|p_{k}^{i}-\pi_{k}\left(T_{k} X_{k-1}^{i}\right)\right\|^{2},
$$

and all the PnP variants that we presented were closed-form methods for minimizing the reprojection error.
Another way to utilize more than 3 points, would be to minimize directly the above cost function using a non-linear optimization approach such as Gauss-Newton or Levenberg-Marquardt.
Under the assumption that errors are zero-mean Gaussian random vector, these solvers return the optimal solution to this non-linear least squares problem. However, they are sensitive to initial conditions and usually solution obtained by P3P or EPnP is used as the initial guess.

## VO and PnP

How is this approach relevant for visual odometry? The stereo images from $k-1$ are used to reconstruct the 3D points in the camera frame $C_{k-1}$ that acts as the world frame from the previous slides. Correspondences are then found in the image $I_{l, k}$ or $I_{r, k}$ and the camera frame $C_{k}$ is aligned to $C_{k-1}$ using a PnP method to obtain the final transformation $T_{k}$.


## Stereo 3D-2D odometry

The pipeline of a 3D-2D odometry might have the following structure:

1. Capture two stereo image pairs $I_{l, k-1}, I_{r, k-1}$ and $I_{l, k}, I_{r, k}$.
2. Detect and match features between $I_{l, k}$ and $I_{l, k-1}$.
3. Match and triangulate matched features in $I_{l, k-1}, I_{r, k-1}$
4. Compute the relative transformation $T_{k}$ using 3D-2D method, e.g., P3P from the triangulated 3D point features $X_{k-1}$ and their reprojections $p_{k}$
5. Concatenate camera transformation by computing $C_{k}=C_{k-1} T_{k}$.

This pipeline models stereo 3D-2D odometry, the monocular case is discussed in the next slide.

The pipeline of a 3D-2D odometry might have the following structure:

1. Do only once:
2. Capture two image frames $I_{k-2}, I_{k-1}$.
3. Detect and match features between them.
4. Triangulate matched features from
5. Do at each iteration:
6. Capture new frame $I_{k}$.
7. Extract and match with previous frame $I_{k-1}$.
8. Compute camera pose (PnP) from 3D-2D matches $C_{k}$.
9. Triangulate new feature matches between $I_{k}, I_{k-1}$.
10. Concatenate transformation by computing $C_{k}=C_{k-1} T_{k}$.

## Outline

- Problem formulation
- 3D-3D motion estimation
-3D-2D motion estimation
- Perspective from $n$ points (PnP)
- Stereo and mono odometry using PnP
- 2D-2D motion estimation
- Estimating the essential matrix
- Relative scale
- Keyframe selection


## 2D-2D motion estimation

The 2D-2D approach estimates motion from 2D-to-2D image point feature correspondences - $p_{k-1}$ and $p_{k}$.
The minimal-case solution involves 5 points.
In general, the goal is to estimate the essential matrix from 2D point correspondences

$$
E_{k} \simeq\left[t_{k}\right]_{\times} R_{k}
$$

Where $\left[t_{k}\right]_{\times}$is the skew-symmetric matrix representation of the translation vector and symbol $\simeq$ means that the equivalence is valid up to a multiplicative scalar.
In other words, we can determine translation only up to scale, i.e., we have a unit length translation vector pointing in the direction of motion.

## Essential matrix via 5 points

The main property of 2D-2D motion estimation methods is the epipolar constraint, which determines the line on which the corresponding feature point $p_{k}^{i}$ lies in the other image.
The minimal case solution involves five 2D-2D correspondences and an efficient implementation is called the Nister's algorithm ${ }^{10}$.
It has become the standard for 2D-2D
 motion estimation in the presence of outliers.

## Essential matrix via 8 points

A simple and straightforward solution for $n \geq 8$ noncoplanar points is the Longuet-Higgins' eight-point algorithm ${ }^{11}$ that we covered in the multiple view geometry part of the course.

Eight-point algorithm - degenerate when the 3D points are coplanar, works for both calibrated and uncalibrated cameras.

Five-point algorithm - works also for
 coplanar points but assumes that the camera is calibrated.

## Relative scale of the translation

Once we have obtained the essential matrix, we can obtain the relative rotation and translation $R_{k}, \hat{t}_{k}$. To recover the trajectory of an image sequence, recall that the different transformations $T_{0: k}$ have to be concatenated.
Since translation is a unit vector, by using this approach frame-by-fame we would obtain a camera trajectory where all the relative translation vectors would be of unit length.


However, not all translations are of the same magnitude, which begs the question if it is possible to take this into account within a 2D-2D framework?

## Relative scale of the translation

This brings us to the notion of the relative scale.
Common approach to determinig the relative scale is to triangulate 3D points $X_{k-1, k}$ and $X_{k, k+1}$ from two subsequent image pairs and the relative distances between any combination of two 3D points can be computed.
The proper scale can then be determined from the distance ratio $r$ between a point pair in $X_{k-1, k}$ and a pair in $X_{k, k+1}$

$$
r_{k}=\frac{\left\|X_{k-1, k}^{i}-X_{k-1, k}^{j}\right\|}{\left\|X_{k, k+1}^{i}-X_{k, k+1}^{j}\right\|}
$$

For robustness, the scale ratios for many point pairs are computed and the mean is used. The translation vector is then scaled with this distance ratio requires features to be matched over multiple frames (at least three).

## Relative scale of the translation

We bootstrap the monocular odometry by computing the relative transformation between the first two views via essential matrix decomposition (translation norm is unit). Then, we triangulate points from the first two views, and as the third image is captured, we triangulate from the second and third view.

The ratio of distance norms between image pairs gets us the relative scale w.r.t to the first pair etc.


## 2D-2D odometry pipeline

The pipeline of a 2D-2D odometry might have the following structure:

1. Capture new frame $I_{k}$.
2. Detect and match features between $I_{k}$ and $I_{k-1}$.
3. Compute essential matrix for image pair $I_{k-1}, I_{k}$
4. Decompose essential matrix into $R_{k}, t_{k}$ and form $T_{k}$.
5. Compute relative scale and rescale $t_{k}$ accordingly .
6. Concatenate transformation by computing $C_{k}=C_{k-1} T_{k}$.

Mmotion estimation based on $E_{k} \simeq\left[t_{k}\right]_{\times} R_{k}$, cannot handle pure rotation, there must exists at least some translational component in the overall motion (to handle this homography estimation is classically used).

## Feature matches - camera shift with frontoparallel scene

feature matches depending on the
camera displacement

translation to the right

rotation to the right
rotation around the optical


## Outline

- Problem formulation
- 3D-3D motion estimation
-3D-2D motion estimation
- Perspective from n points (PnP)
- Stereo and mono odometry using PnP
- 2D-2D motion estimation
- Estimating the essential matrix
- Relative scale
- Keyframe selection

Some of the previous motion estimation methods require triangulation of 3D points.
Triangulated 3D points are determined by intersecting backprojected rays from 2D image correspondences of at least two image frames. In reality, more than 2 lines never intersect in a single point due to:

- image noise
- camera model and calibration errors
- and feature matching uncertainty.

The point at minimal distance from all intersecting rays can be taken as an estimate of the 3D point position or a least squares solution can be found.


When frames are taken at relatively close positions with respect to scene depth (i.e., baseline is much smaller than the feature depth), 3D triangulated points will exhibit large depth uncertainty.
This can affect motion estimation accuracy and in general features should be triangulated at good baseline to depth ratios.


One way to avoid this consists of skipping frames until the average uncertainty of the 3D points decreases below a certain threshold. The selected frames are called keyframes.
Keyframe selection is a very important step in VO and should always be done before updating the motion.
Rule of the thumb: add a keyframe when the following threshold is met (e.g., $10-20 \%$ )

$$
\frac{\text { keyframe_distance }}{\text { average_depth }}>\tau \text {. }
$$



## Questions?


[^0]:    ${ }^{1}$ A. S. Arun, T. S. Huang and S. D. Blostein (1987). „Least-Squares Fitting of Two 3-D Point Sets"

[^1]:    ${ }^{5}$ J. A. Grunert (1841). „Das Pothenotische Problem in erweiterter Gestalt nebst Über seine Anwendungen in der Geodäsie."
    ${ }^{6}$ Gao, Hou, Tang, Cheng (2003). „Complete Solution Classification for the Perspective-Three-Point Problem."

[^2]:    ${ }^{7}$ C. B. Garcia, T. Y. Li (1980). „On the Number of Solutions to Polynomial Systems of Equations."

[^3]:    ${ }^{9}$ V. Lepetit, F. Moreno-Noguer, P. Fua (2009). „EPnP: An Accurate O(n) Solution to the PnP Problem."

